

---

# **LWA Correlator**

***Release 1.0.0***

**Jack Hickish**

**Dec 07, 2020**



## CONTENTS:

<b>1</b>	<b>Installation</b>	<b>1</b>
1.1	Get the Source Code . . . . .	1
1.2	Install Prerequisites . . . . .	1
1.3	Install the Pipeline . . . . .	3
<b>2</b>	<b>Introduction</b>	<b>5</b>
<b>3</b>	<b>Hardware</b>	<b>7</b>
<b>4</b>	<b>Pipeline</b>	<b>9</b>
<b>5</b>	<b>Software</b>	<b>11</b>
5.1	High-Level Parameters . . . . .	11
5.2	Bifrost Block Description . . . . .	11
<b>6</b>	<b>Output Data Formats</b>	<b>15</b>
6.1	Full Correlation Packets . . . . .	15
6.2	Partial Correlation Packets . . . . .	16
6.3	VLBI Beam . . . . .	17
6.4	Integrated Beams . . . . .	17
<b>7</b>	<b>Control Interface</b>	<b>19</b>
7.1	Capture Thread (blockname: capture) . . . . .	19
7.2	Copy Thread (blockname: gpucopy) . . . . .	20
7.3	Correlation Thread (blockname: corr) . . . . .	20
7.4	Visibility Sub-Select Thread (blockname: corrsubsel) . . . . .	21
7.5	Visibility Integrator (blockname: corrace) . . . . .	22
7.6	Beamformer (blockname: beamform) . . . . .	23
<b>8</b>	<b>Indices and tables</b>	<b>25</b>
	<b>Index</b>	<b>27</b>



## INSTALLATION

The LWA 352 Correlator/Beamformer pipeline is available at <https://github.com/realtime-radio/caltech-bifrost-dsp>. Follow the following instructions to download and install the pipeline.

Specify the build directory by defining the `BUILDDIR` environment variable, eg:

```
export BUILDDIR=~/.src/  
mkdir -p $BUILDDIR
```

### 1.1 Get the Source Code

Clone the repository and its dependencies with:

```
# Clone the main repository  
cd $BUILDDIR  
git clone https://github.com/realtime-radio/caltech-bifrost-dsp  
# Clone relevant submodules  
cd caltech-bifrost-dsp  
git submodule init  
git submodule update
```

### 1.2 Install Prerequisites

The following libraries should be installed via the Ubuntu package manager:

```
apt install exuberant-ctags build-essential autoconf libtool exuberant-ctags libhwloc-  
↳dev python3-venv
```

The following 3rd party libraries must also be obtained and installed:

## 1.2.1 CUDA

CUDA can be installed as follows:

```
# Make a directory for the cuda source
mkdir -p $BUILDDIR/cuda
cd $BUILDDIR/cuda
# Download the CUDA installer
wget http://developer.download.nvidia.com/compute/cuda/11.0.2/local_installers/cuda_
↪11.0.2_450.51.05_linux.run

# blacklist nouveau drivers before installing nvidia drivers
sudo su
echo "blacklist nouveau" > /etc/modprobe.d/blacklist-nouveau.conf
echo "options nouveau modeset=0" >> /etc/modprobe.d/blacklist-nouveau.conf
update-initramfs -u
exit

# reboot the machine and nouveau drivers [hopefully] won't start
# reboot now
```

After rebooting, install the CUDA libraries

```
cd $BUILDDIR/cuda
sudo sh cuda_11.0.2_450.51.05_linux.run
# Add CUDA executables to $PATH
echo "export PATH=/usr/local/cuda/bin:${PATH}" >> ~/.bashrc
source ~/.bashrc
```

This CUDA install script will take a minute to unzip and run the installer. If it fails, log messages can be found in `/var/log/nvidia-installer.log` and `/var/log/cuda-installer.log`.

## 1.2.2 IB Verbs

The LWA pipeline uses Infiniband Verbs for fast UDP packet capture. The recommended version is 4.9 LTS. This can be obtained from [https://www.mellanox.com/support/mlnx-ofed-matrix?mtag=linux\\_sw\\_drivers](https://www.mellanox.com/support/mlnx-ofed-matrix?mtag=linux_sw_drivers)

## 1.2.3 xGPU

xGPU is submoduled in the main pipeline repository, to ensure version compatibility. Install with:

```
cd $BUILDDIR/caltech-bifrost-dsp
./install_xgpu
```

## 1.2.4 Bifrost

Bifrost is submoduled in the main pipeline repository, to ensure version compatibility.

The version provided requires Python  $\geq 3.5$ . It is recommended that the bifrost package is installed within a Python version environment.

To install bifrost:

```
cd $BUILDDIR/caltech-bifrost-dsp/bifrost
make
make install
```

## 1.3 Install the Pipeline

After installing the prerequisites above, the LWA pipeline can be installed with

```
cd $BUILDDIR/caltech-lwa-dsp/pipeline
# Be sure to run the installation in the
# appropriate python environment!
python setup.py install
```

---

**Author** Jack Hickish





## **INTRODUCTION**

The LWA-352 X-Engine processing system performs correlation, beamforming, and triggered voltage recording for the LWA-352 array. This document outlines the hardware (Section [sec:hardware]) and software (Section [sec:software]) which makes up the X-Engine, and details the user control interface (Section [sec:api]).



## HARDWARE

The LWA-352 X-engine system comprises 9 1U dual-socket Silicon Mechanics *Rackform R353.v7* servers, each hosting a pair of Nvidia GPUs and solid state memory buffers. Hardware specifications are given in Table [tab:hardware].

Hardware	Model	Notes
Server	Supermicro 1029GQ-TRT	Re-branded as Silicon Mechanics Rackform R353.v7
Motherboard	Supermicro X11 DCQ	
CPU	dual Intel Xeon Scalable Silver 4210R	2.4 GHz, 10 core, 100W TDP
RAM	768 GB PC4-23400	12 x 64 GB; 2933 MHz DDR4; ECC RDIMM
NIC	Mellanox MCX515A-GCAT	ConnectX-5 EN MCX515A-GCAT (1x QSFP28); PCIe 3.0x16
NVMe Controllers	Asus Hyper M.2 X16 Card V2	2 cards per server
NVMe Memory	8TB Samsung 970 Evo Plus	8 x 1 TB
GPU	Nvidia RTX 2080Ti	2 cards per server



## PIPELINE

The pipeline is launched using the `lwa352-pipeline.py` script. This has the following options:

```
usage: lwa352-pipeline.py [-h] [-f] [-c CONFIGFILE] [-l LOGFILE] [-v]
                        [--fakesource] [--nodata] [--testdatain TESTDATAIN]
                        [--testdatacorr TESTDATACORR]
                        [--testdatacorr_acc_len TESTDATACORR_ACC_LEN]
                        [-a CORR_ACC_LEN] [--nocorr] [--nobeamform]
                        [--nogpu] [--ibverbs] [-G GPU] [-P PIPELINEID]
                        [-C CORES] [-q] [--testcorr] [--useetcd]
                        [--etcdhost ETCDHOST] [--ip IP]
                        [--bufgbytes BUFGBYTES]
                        [--target_throughput TARGET_THROUGHPUT]

LWA-SV ADP DRX Service

optional arguments:
  -h, --help                show this help message and exit
  -f, --fork                Fork and run in the background (default: False)
  -c CONFIGFILE, --configfile CONFIGFILE
                          Specify config file (default: adp_config.json)
  -l LOGFILE, --logfile LOGFILE
                          Specify log file (default: None)
  -v, --verbose            Increase verbosity (default: 0)
  --fakesource             Use a dummy source for testing (default: False)
  --nodata                 Don't generate data in the dummy source (faster)
                          (default: False)
  --testdatain TESTDATAIN
                          Path to input test data file (default: None)
  --testdatacorr TESTDATACORR
                          Path to correlator output test data file (default:
                          None)
  --testdatacorr_acc_len TESTDATACORR_ACC_LEN
                          Number of accumulations per sample in correlator test
                          data file (default: 2400)
  -a CORR_ACC_LEN, --corr_acc_len CORR_ACC_LEN
                          Number of accumulations to start accumulating in the
                          slow correlator (default: 240000)
  --nocorr                 Don't use correlation threads (default: False)
  --nobeamform            Don't use beamforming threads (default: False)
  --nogpu                 Don't use any GPU threads (default: False)
  --ibverbs               Use IB verbs for packet capture (default: False)
  -G GPU, --gpu GPU       Which GPU device to use (default: 0)
  -P PIPELINEID, --pipelineid PIPELINEID
                          Pipeline ID. Useful if you are running multiple
```

(continues on next page)

(continued from previous page)

```
pipelines on a single machine (default: 0)
-C CORES, --cores CORES      Comma-separated list of CPU cores to use (default:
                                0,1,2,3,4,5,6,7)
-q, --quiet                  Decrease verbosity (default: 0)
--testcorr                   Compare the GPU correlation with CPU. SLOW!! (default:
                                False)
--useetcd                    Use etcd control/monitoring server (default: False)
--etcdhost ETCDHOST          Host serving etcd functionality (default: etcdhost)
--ip IP                      IP address to which to bind (default: 100.100.100.101)
--bufgbytes BUFGBYTES       Number of GBytes to buffer for transient buffering
                                (default: 4)
--target_throughput TARGET_THROUGHPUT
                                Target throughput when using --fakesource (default:
                                1000.0)
```

## SOFTWARE

The LWA-352 pipeline comprises ?? independent processes, briefly described below.

1. `capture`: Receive F-engine packets and correctly arrange in buffers for downstream processing. Monitor and record missing packets and network performance statistics.
2. `gpubcopy`: Transfer blocks of data from CPU to GPU, for high-performance computation.
3. `triggered_dump`: Buffer large quantities of time-domain data for triggered dump to disk.
4. `corr`: Correlate data using the `xGPU` library.
5. `corr_output_full`: Output full, accumulated visibility matrices.
6. `corrsubsel`: Down-select a sub-set of the complete visibility matrices.
7. `corr_output_part`: Output subselected visibilities as UDP/IP streams.
8. `beamform`: Form multiple voltage beams.
9. `beamform_vlbi_output`: Package and transmit multiple voltage beams for VLBI purposes.
10. `beamform\_sum_beams`: Form integrated power-spectra for multiple beams.
11. `beamform\_output`: Output accumulated power beams.

### 5.1 High-Level Parameters

- `GSize`: “Gulp size” – the number of samples processed per batch by a processing block.

### 5.2 Bifrost Block Description

The bifrost pipelining framework divides streams of data into *sequences*, each of which has a header describing the stream. Different processing blocks act to perform operations on streams, and may modify sequences and their headers.

Here we summarize the bifrost blocks in the LWA-352 pipeline, and the headers they require (for input sequences) and provide (for output sequences).

## 5.2.1 capture

```
class lwa352_pipeline.blocks.capture_block.Capture (log, fs_hz=196000000,
chan_bw_hz=23925.78125,
input_to_ant=None, nstand=352,
npol=2, *args, **kwargs)
```

### Functionality

This block receives UDP/IP data from an Ethernet network and writes it to a bifrost memory buffer.

### New Sequence Condition

This block starts a new sequence each time the incoming packet stream timestamp changes in an unexpected way. For example, if a large block of timestamps are missed a new sequence will be started. Or, if the incoming timestamps decrease (which might happen if the upstream transmitters are reset) a new sequence is started.

### Input Header Requirements

This block is a bifrost source, and thus has no input header requirements.

### Output Headers

Output header fields are as follows:

Field	Format	Units	Description
time_tag	int		Arbitrary integer, incremented with each new sequence.
sync_time	int	UNIX seconds	Synchronization time (corresponding to spectrum sequence number 0)
seq0	int		Spectra number for the first sample in this sequence
chan0	int		Channel index of the first channel in this sequence
nchan	int		Number of channels in the sequence
fs_hz	double	Hz	Sampling frequency of ADCs
sfreq	double	Hz	Center frequency of first channel in the sequence
bw_hz	int	Hz	Bandwidth of the sequence
nstand	int		Number of stands (antennas) in the sequence
npol	int		Number of polarizations per stand in the sequence
complex	bool		True if the data are complex, False otherwise
nbit	int		Number of bits per sample (or per real/imag part if the samples are complex)
input_to_ant	list[int]		List of input to stand/pol mappings with dimensions [nstand x npol, 2]. E.g. if entry N of this list has value [S, P] then the N-th correlator input is stand S, polarization P.
ant_to_input	list[ints]		List of stand/pol to correlator input number mappings with dimensions [nstand, npol]. E.g. if entry [S, P] of this list has value N then stand S, polarization P of the array is the N-th correlator input

### Data Buffers

*Input data buffer:* None

*Output data buffer:* Complex 4-bit data with dimensions (slowest to fastest) Time x Freq x Stand x Polarization

### Instantiation

#### Parameters



- **log** (Python *logging* object.) – Logging object to which runtime messages should be emitted.
- **fs\_hz** (*int*) – Sampling frequency, in Hz, of the upstream ADCs
- **chan\_bw\_hz** (*float*) – Bandwidth of a single frequency channel in Hz.
- **nstand** (*int*) – Number of stands in the array.
- **npol** (*int*) – Number of polarizations per antenna stand.
- **input\_to\_ant** – An map of correlator input to station / polarization. Provided as an `[nstand x npol, 2]` array such that if `input_to_ant[i] == [S,P]` then the *i*-th correlator input is stand *S*, polarization *P*.

#### Keyword Arguments

##### Parameters

- **fmt** (*string*) – The string identifier of the packet format to be received. E.g “snap2”.
- **sock** (*bifrost.udp\_socket.UDPSocket*) – Input UDP socket on which to receive.
- **ring** (*bifrost.ring.Ring*) – bifrost output data ring
- **core** (*int*) – CPU core to which this block should be bound.
- **nsrc** (*int*) – Number of packet sources. This might mean the number of boards transmitting packets, or, in the case that it takes multiple packets from each board to send a complete set of data, this could be a multiple of the number of source boards.
- **src0** (*int*) – The first source to transmit to this block.
- **max\_payload\_size** (*int*) – The maximum payload size, in bytes, of the UDP packets to be received.
- **buffer\_ntime** (*int*) – The number of time samples to be buffered into the output data ring buffer before it is marked full.
- **utc\_start** (*datetime.datetime*) – ?The time at which the block should begin receiving. Set to `datetime.datetime.now()` to start immediately.
- **ibverbs** (*Bool*) – Boolean parameter which, if true, will cause this block to use an Infiniband Verbs packet receiver. If false, or not provided, a standard UDP socket will be used.

5.2.2 gpucopy

5.2.3 triggered\_dump

5.2.4 corr

5.2.5 corr\_output\_full

5.2.6 corrsubsel

5.2.7 corr\_output\_part

5.2.8 beamform

5.2.9 beamform\_vlbi\_output

5.2.10 beamform\_sum\_beams

5.2.11 beamform\_output

## OUTPUT DATA FORMATS

This section defines the output packet formats for each of the pipeline output data products. Unless otherwise specified, all data products are transmitted in network- (i.e. big-) endian format.

Packet sizing is partially determined by the pipeline configuration. Specifically:

- NCHAN – The number of channels processed per pipeline.
- NSTAND – The number of antenna stands in the array.
- NPOL – The number of polarizations per antenna stand.

For the LWA-352 system:

- NCHAN = 184
- NSTAND = 352
- NPOL = 2

### 6.1 Full Correlation Packets

Data from the full, slow correlator are transmitted as a series of UDP packets, with each packet carrying data for one dual-polarization baseline, for multiple channels. Each packet has a 56 byte header followed by a payload of signed 32-bit integers.

```
struct corr_output_full_packet {  
    uint64_t    sync_time;  
    uint64_t    spectra_id;  
    double      bw_hz;  
    double      sfreq_hz;  
    uint32_t    acc_len;  
    uint32_t    nchans;  
    uint32_t    chan0;  
    uint32_t    npols;  
    uint32_t    stand0;  
    uint32_t    stand1;  
    int32_t     data[npols, npols, nchans, 2];  
};
```

Packet fields are as follows:

Field	Format	Units	Description
sync_time	uint64	UNIX sec- onds	The sync time to which spectra IDs are referenced.
spectra_id	int		The spectrum number for the first spectra which contributed to this packet's integration.
bw_hz	double (bi- nary64)	Hz	The total bandwidth of data in this packet
sfreq_hz	double (bi- nary64)	Hz	The center frequency of the first channel of data in this packet
acc_len	uint32		The number of spectra integrated in this packet
nchans	uint32		The number of frequency channels in this packet. For LWA-352 this is 184
chan0	uint32		The index of the first frequency channel in this packet
npols	uint32		The number of polarizations of data in this packet. For LWA-352, this is 2.
stand0	uint32		The index of the first antenna stand in this packet's visibility.
stand1	uint32		The index of the second antenna stand in this packet's visibility.
data	int32*		The data payload. Data for the visibility of antennas at stand0 and stand1, with stand1 conjugated. Data are a multidimensional array of 32-bit integers, with dimensions [NPOLS, NPOLS, NCHANS, 2]. The first axis is the polarization of the antenna at stand0. The second axis is the polarization of the antenna at stand1. The third axis is frequency channel. The fourth axis is complexity, with index 0 the real part of the visibility, and index 1 the imaginary part.

## 6.2 Partial Correlation Packets

Data from the fast dump correlator are transmitted as a series of UDP packets, with each packet carrying data for multiple frequency channels of multiple, single-polarization visibilities.

Each packet has a variable length header followed by a payload of signed 32-bit integers.

```

struct corr_output_partial_packet {
    uint64_t   sync_time;
    uint64_t   spectra_id;
    double     bw_hz;
    double     sfreq_hz;
    uint32_t   acc_len;
    uint32_t   nvis;
    uint32_t   nchans;
    uint32_t   chan0;
    uint32_t   baselines[nvis, 2, 2];
    int32_t    data[nvis, nchans, 2];

```

Packet fields are as follows:

Field	Format	Units	Description
sync_time	uint64	UNIX sec-onds	The sync time to which spectra IDs are referenced.
spectra_id	int		The spectrum number for the first spectra which contributed to this packet's integration.
bw_hz	double (binary64)	Hz	The total bandwidth of data in this packet
sfreq_hz	double (binary64)	Hz	The center frequency of the first channel of data in this packet
acc_len	uint32		The number of spectra integrated in this packet
nvis	uint32		The number of single polarization visibilities present in this packet.
nchans	uint32		The number of frequency channels in this packet. For LWA-352 this is 184
chan0	uint32		The index of the first frequency channel in this packet
baselines	uint32*		An array containing the stand and polarization indices of the multiple visibilities present in this packet. This entry has dimensions [nvis, 2, 2]. The first index runs over the number of visibilities within this packet. The second index is 0 for the first (unconjugated) visibility input and 1 for the second (conjugated) antenna input. The third index is zero for stand number, and 1 for polarization number.
data	int32*		The data payload. Data for the visibility of antennas at stand0 and stand1, with stand1 conjugated. Data are a multidimensional array of 32-bit integers, with dimensions [NVIS, NCHANS, 2]. The first axis runs over the multiple visibilities in this packet. Each index can be associated with a physical antenna using the <code>baselines</code> field. The second axis is frequency channel. The third axis is complexity, with index 0 the real part of the visibility, and index 1 the imaginary part.

## 6.3 VLBI Beam

## 6.4 Integrated Beams



## CONTROL INTERFACE

Control and monitoring of the X-Engine pipeline is carried out through the passing of JSON-encoded messages through an `etcd`<sup>1</sup> key-value store. Each processing block in the LWA system has a unique identifier which defines a key to which runtime status is published and a key which should be monitored for command messages.

The unique key of a processing block is derived from the `blockname` of the module within the pipeline, the `hostname` of the server on which a pipeline is running, and the pipeline id - `pid` - of this pipeline.

In general, keys to which status information is published have the prefix:

```
/mon/corr/xeng/<hostname>/pipeline/<pid>/<blockname>.
```

Keys to which users should write commands have the prefix

```
/cmd/corr/xeng/<hostname>/pipeline/<pid>/<blockname>.
```

The format of these status and command messages, and their allowed values are given in the remainder of this section on a per-block basis.

### 7.1 Capture Thread (blockname: `capture`)

#### 7.1.1 Commands

The `capture` block accepts no runtime commands. When a pipeline is executed, the capture module will automatically begin filling processing buffers. Buffer boundaries occur every `G_SIZE` samples.

#### 7.1.2 Monitoring

The `capture` block writes monitoring data to the key `/mon/corr/xeng/<hostname>/pipeline/<pid>/capture`. Data are written as a JSON-encoded dictionary with the following entries:

Field	Format	Units	Description
<code>throughput</code>	float	Gbits/s	Block throughput
<code>n_dropped</code>	int	packets	Number of packets dropped since pipeline start
<code>n_received</code>	int	packets	Number of packets received since pipeline start
<code>frac_dropped</code>	float		Fraction of packets dropped since pipeline start
<code>n_late</code>	int	packets	Number of late packets since pipeline start
<code>n_f_missing</code>	int	boards	TODO
<code>n_part_dropped</code>	int	packets	TODO
<code>time</code>	float	UNIX time	The time this key was updated.

---

<sup>1</sup> See `etcd.io`

## 7.2 Copy Thread (blockname: `gpubcopy`)

The `gpubcopy` block accepts no runtime commands and outputs no run-time statistics.

## 7.3 Correlation Thread (blockname: `corr`)

The `corr` block takes blocks of `G_SIZE` 4-bit time samples from the `gpubcopy` thread and generates visibility matrices using an xGPU computation kernel. Integration takes place over the `G_SIZE` input samples.

### 7.3.1 Commands

The `corr` block has a run-time configurable accumulation length and start time. These can be set by writing a JSON-encoded dictionary to the key `/cmd/corr/xeng/<hostname>/pipeline/<pid>/corr`, which should have the following fields:

`c c c X` Field & Format & Units & Description

`acc_len` & int & samples & Number of samples to integrate. Must be a multiple of `G_SIZE`. `acc_len = 0` can be used to force the `corr` module to stop processing.

`start_time` & int & samples & Sample index on which to begin integrating. Must be a multiple of `G_SIZE`.

Sample indices are relative to the F-Engine sync time – i.e., sample index 0 is the first sample after an F-Engine sync event. Sample indices can only be converted to real time with the knowledge of the F-Engine sync time and F-Engine ADC clock rate.

It should be noted that modifying the run-time configuration of the `corr` module will impact both the fast- and slow-visibility processing streams. Both streams will re-synchronize onto new correlator integration boundaries.

### 7.3.2 Monitoring

The `corr` block writes status data as a JSON-encoded dictionary to the key:

`/mon/corr/xeng/<hostname>/pipeline/<pid>/corr`

The status dictionary has the following fields:

`c c c X` Field & Format & Units & Description

`throughput` & float & Gbits/s & Block throughput

`acc_len` & int & samples & Number of samples currently set to integrate

`start_sample` & int & samples & Current start time.

`curr_sample` & int & samples & The last sample to be processed.

`update_pending` & bool & - & True if new integration parameters are waiting to be loaded.

`last_update_time` & float & seconds & The time since UNIX epoch that the integration parameters were last updated.

`new_acc_len` & int & samples & The commanded integration length

`new_start_sample` & int & samples & The commanded start sample

`last_cmd_time` & float & seconds & The time since UNIX epoch that the last command was received



## 7.4 Visibility Sub-Select Thread (blockname: `corrsubsel`)

### 7.4.1 Commands

The `corrsubsel` block outputs a run-time configurable set of baselines. These can be set by writing a JSON-encoded dictionary to the key `/cmd/corr/xeng/<hostname>/pipeline/<pid>/corrsubsel`, which should have the following fields:

`c c c X` Field & Format & Units & Description

`subsel` & `list(int)` & - & A list of baselines for subselection. This field should be provided as a multidimensional list with dimensions `[N_VIS, 2, 2]`. The first axis runs over the 4656 baselines which may be selected. The second index is 0 for the first (unconjugated) input selected and 1 for the second (conjugated) input selected. The third axis is 0 for stand number, and 1 for polarization number.

### Example

To set the baseline subsection to choose:

- visibility 0: the autocorrelation of antenna 0, polarization 0
- visibility 1: the cross correlation of antenna 5, polarization 1 with antenna 6, polarization 0

use:

```
subsel = [ [[0,0], [0,0]], [[5,1], [6,0]], ... ]
```

Note that the uploaded selection list must always have 4656 entries.

### 7.4.2 Monitoring

The `corr` block writes status data as a JSON-encoded dictionary to the key: `/mon/corr/xeng/<hostname>/pipeline/<pid>/corrsubsel`.

The status dictionary has the following fields:

Field	Format	Units	Description
<code>throughput</code>	float	Gbits/s	Block throughput
<code>subsel</code>	<code>list(int)</code>	samples	Current set of visibility indices being selected
<code>update_pending</code>	bool		True if new selection parameters are waiting to be loaded.
<code>last_update_time</code>	float	seconds	The time since UNIX epoch that the selection parameters were last updated.
<code>new_subsel</code>	<code>list(int)</code>	samples	The commanded visibility selection indices.
<code>last_cmd_time</code>	float	seconds	The time since UNIX epoch that the last command was received

## 7.5 Visibility Integrator (blockname: `corracc`)

### 7.5.1 Commands

The `corracc` block further integrates the output of the `corr` block. Integration parameters can be set by writing a JSON-encoded dictionary to the key:

```
/cmd/corr/xeng/<hostname>/pipeline/<pid>/corracc
```

This should have the following fields:

Field	Format	Units	Description
<code>acc_len</code>	int		Number of samples to integrate. <code>acc_len = 0</code> can be used to force the <code>corracc</code> module to stop processing.
<code>start_time</code>	int	samples	Sample index on which to begin integrating.

Note that the `acc_len` configuration must be compatible with – i.e., must be a multiple of – the accumulation length set in the `corr` block. Furtherm the `start_time` must be compatible with the integration boundaries associated with the `corr` block’s integration settings.

Run-time checks will flag bad configurations as errors, but no check is made on issuing a command to ensure it is valid. After booting the pipeline, a safe order of configuration is:

1. Boot pipeline.
2. Configure `corracc` block
3. Configure `corr` block

For changes of configuration, the safe order of updates is:

### 7.5.2 Monitoring

The `corracc` block writes status data as a JSON-encoded dictionary to the key:

```
/mon/corr/xeng/<hostname>/pipeline/<pid>/corracc.
```

The status dictionary has the following fields:

Field	Format	Units	Description
acc_len	int	samples	Number of samples currently set to integrate
start_sample	int	samples	Current start time.
curr_sample	int	samples	The last sample to be processed.
update_pending	bool		True if new integration parameters are waiting to be loaded.
last_update_time	float	seconds	The time since UNIX epoch that the integration parameters were last updated.
new_acc_len	int	samples	The commanded integration length
new_start_sample	int	samples	The commanded start sample
last_cmd_time	float	seconds	The time since UNIX epoch that the last command was received

## 7.6 Beamformer (blockname: beamform)

The `beamform` block forms  $2 \times \text{NBEAM}$  independent, single polarization voltage beams. Beam pointings are specified by relative antenna delays and a set of universal, frequency-dependent calibration coefficients, which are shared among all beams. Note that this interface precludes direction-dependent calibrations.

### 7.6.1 Commands

Commands are sent to the `beamform` module by writing a JSON-encoded command to the key:

```
/cmd/corr/xeng/<hostname>/pipeline/<pid>/beamform
```

This command should have the following fields

Field	Format	Units	Description
delays[x]	list(float)	ns	An <code>NINPUT</code> element list of geometric delays, in nanoseconds. [x] is a beam index, and should be between 0 and <code>NBEAM - 1</code>
gains	list(complex32)		A two dimensional list of calibration gains with shape <code>[NCHAN, NINPUT]</code>
load_sample	int	sample	Sample number on which the supplied delays should be loaded. If this field is absent, new delays will be loaded as soon as possible

The `beamform` block calculates voltage beams only and has no concept of polarization. Instead, the `beamform` block generates  $2 \times \text{NBEAM}$  beams and computes the auto- and cross-power spectra between beams in order to generate auto- and cross-pol products. Beams are paired such that the cross-power of beams  $2n$  and  $2n+1$  are computed – it is the user’s responsibility to ensure that these beams have the same pointing and are formed from complementary antenna polarizations.

## 7.6.2 Monitoring

The `beamform` block writes status data as a JSON-encoded dictionary to the key:

```
/mon/corr/xeng/<hostname>/pipeline/<pid>/beamform.
```

The status dictionary has the following fields:

Field	Format	Units	Description
throughput	float	Gbits/s	Block throughput
delays[x]	list(float)	ns	An <code>NINPUT</code> element list containing the delays currently loaded for beam <code>x</code>
gains	list(complex32)		A two dimensional list of currently loaded calibration gains. The dimensions of this list should be <code>NCHAN × NINPUT</code>
new_delays[x]	list(float)	ns	An <code>NINPUT</code> element list containing the next set of delays to be loaded for beam <code>x</code>
new_gains	list(complex32)		A two-dimensional list of calibration gains with shape <code>[NCHAN, NINPUT]</code>
curr_sample	int	samples	The last sample to be processed.
update_pending	bool	•	True if new integration parameters are waiting to be loaded.
last_update_time	float	seconds	The time since UNIX epoch that the integration parameters were last updated.
new_acc_len	int	samples	The commanded integration length
new_start_sample	int	samples	The commanded start sample
last_cmd_time	float	seconds	The time since UNIX epoch that the last command was received

## INDICES AND TABLES

- `genindex`
- `modindex`
- `search`



## INDEX

### C

Capture (*class in lwa352\_pipeline.blocks.capture\_block*),  
[12](#)